

Project Details Are From Page 09

Ingress

Ingress Networking

Ingress networking in Kubernetes refers to the way external access to services within a Kubernetes cluster is managed. It provides a powerful mechanism to define rules that allow external HTTP and HTTPS traffic to reach your services. An Ingress resource defines these rules and allows for more complex routing, such as load balancing, SSL termination, and name-based virtual hosting.

Key Concepts

- 1 **Ingress Resource:** A set of rules that allow inbound connections to reach the cluster services.
- 2 **Ingress Controller:** A component that implements the Ingress resource, typically deployed as a pod within the cluster. It interprets the Ingress rules and routes traffic accordingly.
- 3 **Ingress Rules:** Define the routing of HTTP/HTTPS traffic to various services within the cluster based on host, path, etc.
- 4 **Backend Services:** Kubernetes services that Ingress resources route traffic to.

1. Start Minikube with Ingress Addon

Ensure that Minikube is started with the Ingress addon enabled.

```
minikube start --addons=ingress
```

2. Verify Ingress Controller

Check that the NGINX Ingress controller pods are running.

```
kubectl get pods -n kube-system | grep nginx
```

You should see the NGINX Ingress controller pods listed.

3. Create Sample Services and Deployments

Create sample deployments and services for demonstration.

```
# frontend-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
```

```
  name: frontend
spec:
  replicas: 2
  selector:
    matchLabels:
      app: frontend
  template:
    metadata:
      labels:
        app: frontend
    spec:
      containers:
        - name: frontend
          image: daradesudarshan/centralrepo:nginx-app-frontend
          ports:
            - containerPort: 80
```

```
---
apiVersion: v1
kind: Service
metadata:
  name: frontend-service
spec:
  selector:
    app: frontend
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

```
# backend-deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: backend
spec:
  replicas: 2
  selector:
    matchLabels:
      app: backend
  template:
    metadata:
      labels:
        app: backend
    spec:
      containers:
        - name: backend
```

```

        image: daradesudarshan/centralrepo:nginx-app-backend
        ports:
          - containerPort: 80
---
apiVersion: v1
kind: Service
metadata:
  name: backend-service
spec:
  selector:
    app: backend
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80

```

Apply these YAML files to create the deployments and services.

```

kubectl apply -f frontend-deployment.yaml
kubectl apply -f backend-deployment.yaml

```

4. Create an Ingress Resource

Define an Ingress resource to route traffic to these services.

```

# ingress-resource.yaml
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example-ingress
spec:
  rules:
    - host: myapp.local
      http:
        paths:
          - path: /frontend
            pathType: Prefix
            backend:
              service:
                name: frontend-service
                port:
                  number: 80
          - path: /backend
            pathType: Prefix

```

```
backend:
  service:
    name: backend-service
    port:
      number: 80
```

Apply the Ingress resource.

```
kubectl apply -f ingress-resource.yaml
```

5. Update `/etc/hosts`

Add the hostname defined in the Ingress resource to your `/etc/hosts` file pointing to the Minikube IP.

```
sudo nano /etc/hosts
```

Add the following line (replace `<minikube-ip>` with the actual Minikube IP):

```
<minikube-ip> myapp.local
```

Get the Minikube IP using:

```
minikube ip
```

6. Access the Services

You should now be able to access the services via your browser or `curl`.

```
curl http://myapp.local/frontend
curl http://myapp.local/backend
```

Advanced Use Cases

1. Load Balancing with Sticky Sessions

Scenario: Distribute traffic across multiple instances of a service while maintaining session affinity.

Example:

Ingress Resource with Annotations:

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: sticky-ingress
  annotations:
    nginx.ingress.kubernetes.io/affinity: "cookie"
    nginx.ingress.kubernetes.io/session-cookie-name: "route"
spec:
  rules:
  - host: sticky.local
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: sticky-service
            port:
              number: 80
```

2. Path Rewriting

Scenario: Route traffic to a different path in the backend service.

Example:

Ingress Resource with Path Rewriting:

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: rewrite-ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  rules:
  - host: rewrite.local
    http:
      paths:
      - path: /oldpath/(.*)
        pathType: Prefix
        backend:
          service:
            name: new-service
            port:
```

number: 80

nginx.ingress.kubernetes.io/rewrite-target: /

- **Description:** Replaces the path with the specified rewrite target (/). It effectively removes the matched part of the URL path.
- **Example:** If the original request path is `/oldpath/something`, it will be rewritten to `/something` before forwarding to the backend service.

nginx.ingress.kubernetes.io/rewrite-target: /newpath

- **Description:** Replaces the path with the specified rewrite target (/newpath). It substitutes the matched part of the URL path with the specified new path.
- **Example:** If the original request path is `/oldpath/something`, it will be rewritten to `/newpath/something` before forwarding to the backend service.

nginx.ingress.kubernetes.io/rewrite-target: /\$1

- **Description:** Uses a capture group (\$1) from the original path's regular expression match to dynamically construct the rewritten path.
- **Example:** If your Ingress path definition includes regex capturing groups like `path: /oldpath/(.*)`, and `rewrite-target: /$1` is specified, requests to `/oldpath/something` will be rewritten to `/something`.

Key Features of TLS:

- 1 **Encryption:** TLS encrypts data to ensure that it remains private and secure during transmission. This prevents unauthorized parties from intercepting and reading the data.
- 2 **Authentication:** TLS supports server-side and optional client-side authentication using digital certificates. This ensures that the parties involved in the communication are who they claim to be.
- 3 **Compatibility:** TLS is widely supported and used across various applications and services, including web browsers, email clients, instant messaging, and more.

TLS Handshake Process:

TLS communication begins with a handshake process, where the client and server negotiate parameters for secure communication:

- **ClientHello:** The client sends a message containing the TLS version, supported cipher suites, and a random number.
- **ServerHello:** The server responds with its chosen TLS version, cipher suite, and a random number.
- **Certificate Exchange:** The server sends its digital certificate to prove its identity (if required).

- **Key Exchange:** The client and server agree on a shared secret key to be used for symmetric encryption during the session.
- **Finished:** Both parties exchange finished messages to confirm that the handshake was successful and communication can proceed securely.

Usage:

TLS is commonly used to secure HTTP connections (HTTPS), ensuring that sensitive information such as login credentials, payment details, and personal data transmitted over the internet remains confidential and integral.

Create a Secret for TLS Certificate

```
kubectl create secret tls tls-secret --cert=path/to/tls.crt --key=path/to/tls.key
```

Ingress Resource:

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: example-ingress
  annotations:
    nginx.ingress.kubernetes.io/ssl-redirect: "true"
spec:
  tls:
  - hosts:
    - myapp.local
    secretName: my-tls-secret
  rules:
  - host: myapp.local
    http:
      paths:
      - path: /frontend
        pathType: Prefix
        backend:
          service:
            name: frontend-service
            port:
              number: 80
      - path: /backend
```

```
pathType: Prefix
backend:
  service:
    name: backend-service
    port:

    number: 80
```

Project Overview

Participants are required to deploy a simple static web application on a Kubernetes cluster using Minikube, set up advanced ingress networking with URL rewriting and sticky sessions, and configure horizontal pod autoscaling to manage traffic efficiently. The project will be divided into stages, with each stage focusing on specific aspects of Kubernetes ingress, URL rewriting, sticky sessions, and autoscaling.

Requirements and Deliverables

Stage 1: Setting Up the Kubernetes Cluster and Static Web App

1 Set Up Minikube:

- Ensure Minikube is installed and running on the local Ubuntu machine.
- Verify the Kubernetes cluster is functioning correctly.

```
vagrant@ubuntu2204:~$ minikube start
minikube v1.33.1 on Ubuntu 22.04 (vbox/amd64)
Automatically selected the docker driver. Other choices: none, ssh

The requested memory allocation of 1963MiB does not leave room for system overhead (total system memory: 1963MiB). You may face stability issues
Suggestion: Start minikube with less memory allocated: 'minikube start --memory=1963mb'

Using Docker driver with root privileges
Starting "minikube" primary control-plane node in "minikube" cluster
Pulling base image v0.0.44 ...
Creating docker container (CPUs=2, Memory=1963MB) ...
Preparing Kubernetes v1.30.0 on Docker 26.1.1 ...
  ■ Generating certificates and keys ...
  ■ Booting up control plane ...
  ■ Configuring RBAC rules ...
Configuring bridge CNI (Container Networking Interface) ...
Verifying Kubernetes components...
  ■ Using image gcr.io/k8s-minikube/storage-provisioner:v5
Enabled addons: default-storageclass, storage-provisioner
Done! kubectrl is now configured to use "minikube" cluster and "default" namespace by default
vagrant@ubuntu2204:~$
```

2 Deploy Static Web App:

- Create a Dockerfile for a simple static web application (e.g., an HTML page served by Nginx).

```
vagrant@ubuntu2204:~/DevOps-class-assessment$ cat Dockerfile
FROM nginx:latest

COPY . /usr/share/nginx/html

EXPOSE 80
vagrant@ubuntu2204:~/DevOps-class-assessment$
```


- Build a Docker image for the static web application.
- Push the Docker image to Docker Hub or a local registry.

3 Kubernetes Deployment:

- Write a Kubernetes deployment manifest to deploy the static web application.
- Write a Kubernetes service manifest to expose the static web application within the cluster.
- Apply the deployment and service manifests to the Kubernetes cluster.

```
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ kubectl apply -f frontend-deployment.yaml
deployment.apps/frontend created
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ kubectl apply -f backend-deployment.yaml
deployment.apps/backend created
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ kubectl apply -f frontend-service.yaml
service/frontend-service created
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ kubectl apply -f backend-service.yaml
service/backend-service created
```

Deliverables:

- Dockerfile for the static web app
- Docker image URL
- Kubernetes deployment and service YAML files

Stage 2: Configuring Ingress Networking

4 Install and Configure Ingress Controller:

- Install an ingress controller (e.g., Nginx Ingress Controller) in the Minikube cluster.

```
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ minikube addons enable ingress
! ingress is an addon maintained by Kubernetes. For any concerns contact minikube on GitHub.
You can view the list of minikube maintainers at: https://github.com/kubernetes/minikube/blob/master/OWNERS
■ Using image registry.k8s.io/nginx-ingress/controller:v1.10.1
■ Using image registry.k8s.io/nginx-ingress/kube-webhook-certgen:v1.4.1
■ Using image registry.k8s.io/nginx-ingress/kube-webhook-certgen:v1.4.1
🌟 Verifying ingress addon...
🌟 The 'ingress' addon is enabled
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$
```

- Verify the ingress controller is running and accessible.

5 Create Ingress Resource:

- Write an ingress resource manifest to route external traffic to the static web application.
- Configure advanced ingress rules for path-based routing and host-based routing (use at least two different hostnames and paths).

```
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ curl http://myapp.local/frontend
<!DOCTYPE html>
<html>
<body>

<h1>Hello World!</h1>
<p>Welcome to Kubernetes</p>

</body>
</html>
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ curl http://myapp.local/backend
<!DOCTYPE html>
<html>
<body>

<h1>Hello World!</h1>
<p>Welcome to Kubernetes</p>

</body>
</html>
```

- Implement TLS termination for secure connections.

```
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ openssl genrsa -out ca.key 2048
OpenSSL: command not found
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ openssl genrsa -out ca.key 2048
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ ls
backend-deployment.yaml backend-service.yaml ca.key frontend-deployment.yaml frontend-service.yaml ingress-resource.yaml
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ openssl req -x509 \
  -new -nodes \
  -days 365 \
  -key ca.key \
  -out ca.crt \
  -subj "/CN=myapp.local"
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ ls
backend-deployment.yaml backend-service.yaml ca.crt ca.key frontend-deployment.yaml frontend-service.yaml ingress-resource.yaml
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ kubectl create secret tls my-tls-secret \
  --key ca.key \
  --cert ca.crt
secret/my-tls-secret created
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ kubectl get secrets/my-tls-secret
NAME          TYPE          DATA   AGE
my-tls-secret  kubernetes.io/tls  2       13s
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$
```

- Configure URL rewriting in the ingress resource to modify incoming URLs before they reach the backend services.
- Enable sticky sessions to ensure that requests from the same client are directed to the same backend pod.

```
annotations:
  nginx.ingress.kubernetes.io/rewrite-target: /
  nginx.ingress.kubernetes.io/affinity: "cookie"
  nginx.ingress.kubernetes.io/session-cookie-name: "route"
  nginx.ingress.kubernetes.io/ssl-redirect: "true"
```

Deliverables:

- Ingress controller installation commands/scripts
- Ingress resource YAML file with advanced routing, TLS configuration, URL rewriting, and sticky sessions

Stage 3: Implementing Horizontal Pod Autoscaling

6 Configure Horizontal Pod Autoscaler:

- Write a horizontal pod autoscaler (HPA) manifest to automatically scale the static web application pods based on CPU utilization.
- Set thresholds for minimum and maximum pod replicas.

```
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ kubectl autoscale deployment frontend --cpu-percent=50 --min=2 --max=4
horizontalpodautoscaler.autoscaling/frontend autoscaled
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ kubectl autoscale deployment backend --cpu-percent=50 --min=2 --max=4
horizontalpodautoscaler.autoscaling/backend autoscaled
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$
```

7 Stress Testing:

- Perform stress testing to simulate traffic and validate the HPA configuration.
- Monitor the scaling behavior and ensure the application scales up and down based on the load.

Deliverables:

- Horizontal pod autoscaler YAML file
- Documentation or screenshots of the stress testing process and scaling behavior

Stage 4: Final Validation and Cleanup

8 Final Validation:

- Validate the ingress networking, URL rewriting, and sticky sessions configurations by accessing the web application through different hostnames and paths.

```
vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ curl -I https://myapp.local/frontend -Lk
HTTP/2 200
date: Fri, 19 Jul 2024 07:02:40 GMT
content-type: text/html
content-length: 129
set-cookie: route=1721372561.623.26.956366|a3828f09fdccff715aad18c301e466a0; Path=/frontend; Secure; HttpOnly
last-modified: Fri, 19 Jul 2024 02:51:11 GMT
etag: "6699d49f-81"
accept-ranges: bytes

vagrant@ubuntu2204:~/DevOps-class-assessment/assessment10/kubernetes$ curl -I https://myapp.local/backend -Lk
HTTP/2 200
date: Fri, 19 Jul 2024 07:03:04 GMT
content-type: text/html
content-length: 131
set-cookie: route=1721372585.448.25.758690|d529847c6f85bcc10bf781fdaa207d53; Path=/backend; Secure; HttpOnly
last-modified: Fri, 19 Jul 2024 02:52:59 GMT
etag: "6699d50b-83"
accept-ranges: bytes
```

- Verify the application's availability and performance during different load conditions.

Stress test for frontend deployments :

Every 2.0s: kubectl get all -o wide

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
pod/backend-57bb7fb6f-4t9cz	1/1	Running	0	3m42s	10.244.0.28	minikube	<none>	<none>
pod/backend-57bb7fb6f-8kdjj	1/1	Running	0	3m40s	10.244.0.29	minikube	<none>	<none>
pod/frontend-7858c75f57-2l6jx	1/1	Running	0	4m6s	10.244.0.27	minikube	<none>	<none>
pod/frontend-7858c75f57-ftgpn	1/1	Running	0	43s	10.244.0.31	minikube	<none>	<none>
pod/frontend-7858c75f57-mvx85	1/1	Running	0	43s	10.244.0.30	minikube	<none>	<none>
pod/frontend-7858c75f57-pklks	1/1	Running	0	4m7s	10.244.0.26	minikube	<none>	<none>
pod/load-generator	1/1	Running	0	11m	10.244.0.25	minikube	<none>	<none>

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
service/backend-service	ClusterIP	10.96.186.38	<none>	80/TCP	3h5m	app=backend
service/frontend-service	ClusterIP	10.98.109.48	<none>	80/TCP	3h5m	app=frontend
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3h9m	<none>

NAME	READY	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR
deployment.apps/backend	2/2	2	2	3h2m	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend
deployment.apps/frontend	4/4	4	4	3h5m	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR
replicaset.apps/backend-57bb7fb6f	2	2	2	3m42s	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend,pod-template-hash=57bb7fb6f
replicaset.apps/backend-6697d4b966	0	0	0	3h5m	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend,pod-template-hash=6697d4b966
replicaset.apps/backend-bd8549865	0	0	0	34m	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend,pod-template-hash=bd8549865
replicaset.apps/frontend-659fbd7cc	0	0	0	32m	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend,pod-template-hash=659fbd7cc
replicaset.apps/frontend-7858c75f57	4	4	4	4m7s	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend,pod-template-hash=7858c75f57
replicaset.apps/frontend-7cf56cdfd	0	0	0	3h5m	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend,pod-template-hash=7cf56cdfd

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
horizontalpodautoscaler.autoscaling/backend	Deployment/backend	cpu: 0%/50%	2	4	2	68m
horizontalpodautoscaler.autoscaling/frontend	Deployment/frontend	cpu: 100%/50%	2	4	4	68m

Every 2.0s: kubectl get all -o wide

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
pod/backend-57bb7fb6f-2bhnn	1/1	Running	0	9s	10.244.0.33	minikube	<none>	<none>
pod/backend-57bb7fb6f-4t9cz	1/1	Running	0	5m8s	10.244.0.28	minikube	<none>	<none>
pod/backend-57bb7fb6f-8kdjj	1/1	Running	0	5m6s	10.244.0.29	minikube	<none>	<none>
pod/backend-57bb7fb6f-zt6fn	1/1	Running	0	9s	10.244.0.32	minikube	<none>	<none>
pod/frontend-7858c75f57-2l6jx	1/1	Running	0	5m32s	10.244.0.27	minikube	<none>	<none>
pod/frontend-7858c75f57-ftgpn	1/1	Running	0	2m9s	10.244.0.31	minikube	<none>	<none>
pod/frontend-7858c75f57-mvx85	1/1	Running	0	2m9s	10.244.0.30	minikube	<none>	<none>
pod/frontend-7858c75f57-pklks	1/1	Running	0	5m33s	10.244.0.26	minikube	<none>	<none>
pod/load-generator	1/1	Running	0	12m	10.244.0.25	minikube	<none>	<none>

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
service/backend-service	ClusterIP	10.96.186.38	<none>	80/TCP	3h6m	app=backend
service/frontend-service	ClusterIP	10.98.109.48	<none>	80/TCP	3h6m	app=frontend
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3h11m	<none>

NAME	READY	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR
deployment.apps/backend	4/4	4	4	3h7m	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend
deployment.apps/frontend	4/4	4	4	3h7m	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR
replicaset.apps/backend-57bb7fb6f	4	4	4	5m8s	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend,pod-template-hash=57bb7fb6f
replicaset.apps/backend-6697d4b966	0	0	0	3h7m	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend,pod-template-hash=6697d4b966
replicaset.apps/backend-bd8549865	0	0	0	35m	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend,pod-template-hash=bd8549865
replicaset.apps/frontend-659fbd7cc	0	0	0	34m	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend,pod-template-hash=659fbd7cc
replicaset.apps/frontend-7858c75f57	4	4	4	5m33s	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend,pod-template-hash=7858c75f57
replicaset.apps/frontend-7cf56cdfd	0	0	0	3h7m	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend,pod-template-hash=7cf56cdfd

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
horizontalpodautoscaler.autoscaling/backend	Deployment/backend	cpu: 100%/50%	2	4	2	69m
horizontalpodautoscaler.autoscaling/frontend	Deployment/frontend	cpu: 100%/50%	2	4	4	69m

9.Cleanup:

Every 2.0s: kubectl get all -o wide

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
pod/backend-57bb7fb6f-4t9cz	1/1	Running	0	46m	10.244.0.28	minikube	<none>	<none>
pod/backend-57bb7fb6f-8kdjj	1/1	Running	0	46m	10.244.0.29	minikube	<none>	<none>
pod/frontend-7858c75f57-2l6jx	1/1	Running	0	46m	10.244.0.27	minikube	<none>	<none>
pod/frontend-7858c75f57-pklks	1/1	Running	0	46m	10.244.0.26	minikube	<none>	<none>

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
service/backend-service	ClusterIP	10.96.186.38	<none>	80/TCP	3h41m	app=backend
service/frontend-service	ClusterIP	10.98.109.48	<none>	80/TCP	3h41m	app=frontend
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3h46m	<none>

NAME	READY	UP-TO-DATE	AVAILABLE	AGE	CONTAINERS	IMAGES	SELECTOR
deployment.apps/backend	2/2	2	2	3h42m	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend
deployment.apps/frontend	2/2	2	2	3h42m	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend

NAME	DESIRED	CURRENT	READY	AGE	CONTAINERS	IMAGES	SELECTOR
replicaset.apps/backend-57bb7fb6f	2	2	2	46m	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend,pod-template-hash=57bb7fb6f
replicaset.apps/backend-6697d4b966	0	0	0	3h42m	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend,pod-template-hash=6697d4b966
replicaset.apps/backend-bd8549865	0	0	0	71m	backend	daradesudarshan/centralrepo:nginx-app-backend	app=backend,pod-template-hash=bd8549865
replicaset.apps/frontend-659fbd7cc	0	0	0	69m	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend,pod-template-hash=659fbd7cc
replicaset.apps/frontend-7858c75f57	2	2	2	46m	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend,pod-template-hash=7858c75f57
replicaset.apps/frontend-7cf56cdfd	0	0	0	3h42m	frontend	daradesudarshan/centralrepo:nginx-app-frontend	app=frontend,pod-template-hash=7cf56cdfd

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
horizontalpodautoscaler.autoscaling/backend	Deployment/backend	cpu: 0%/50%	2	4	2	104m
horizontalpodautoscaler.autoscaling/frontend	Deployment/frontend	cpu: 0%/50%	2	4	2	104m

- Provide commands or scripts to clean up the Kubernetes resources created during the project (deployments, services, ingress, HPA).

Deliverables:

- Final validation report documenting the testing process and results
- Cleanup commands/scripts